
CUDA, Parallel Programming Language Parallel Patterns

Jiwon Seo

CUDA Common Templates

```
__global__ void foo(...) {  
    extern __shared__ smem[];  
    int i = ???
```

```
    // now what???
```

```
}
```

```
...
```

```
int B = ???
```

```
int N = ???
```

```
int S = ???
```

```
foo<<<B,N,S>>>() ;
```

Parallel Patterns

- Think at a higher level than individual CUDA kernels
- Specify **what** to compute, not **how** to compute it
- Let programmer worry about algorithm
- Defer pattern implementation to someone else

Common Parallel Computing Scenarios

- Many parallel threads need to generate a single result
→ Reduce
- Many parallel threads need to partition data
→ Split
- Many parallel threads produce variable output / thread
→ Compact / Expand

Primordial CUDA Pattern: Blocking

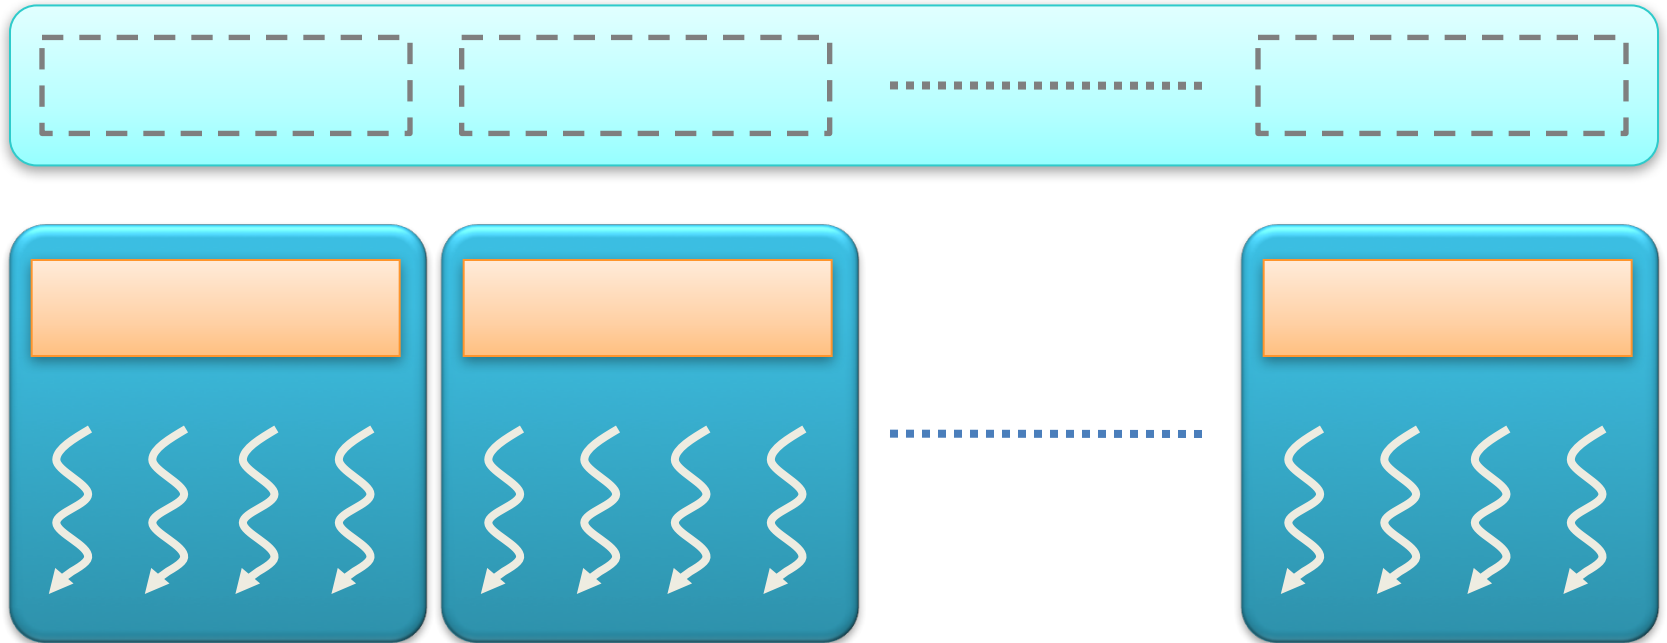
- Partition data to operate in well-sized blocks
 - Small enough to be staged in shared memory
 - Assign each data partition to a thread block
 - No different from cache blocking!
- Provides several performance benefits
 - Have enough blocks to keep processors busy
 - Working in shared memory cuts memory latency dramatically
 - Likely to have coherent access patterns on load/store to shared memory

Primordial CUDA Pattern: Blocking



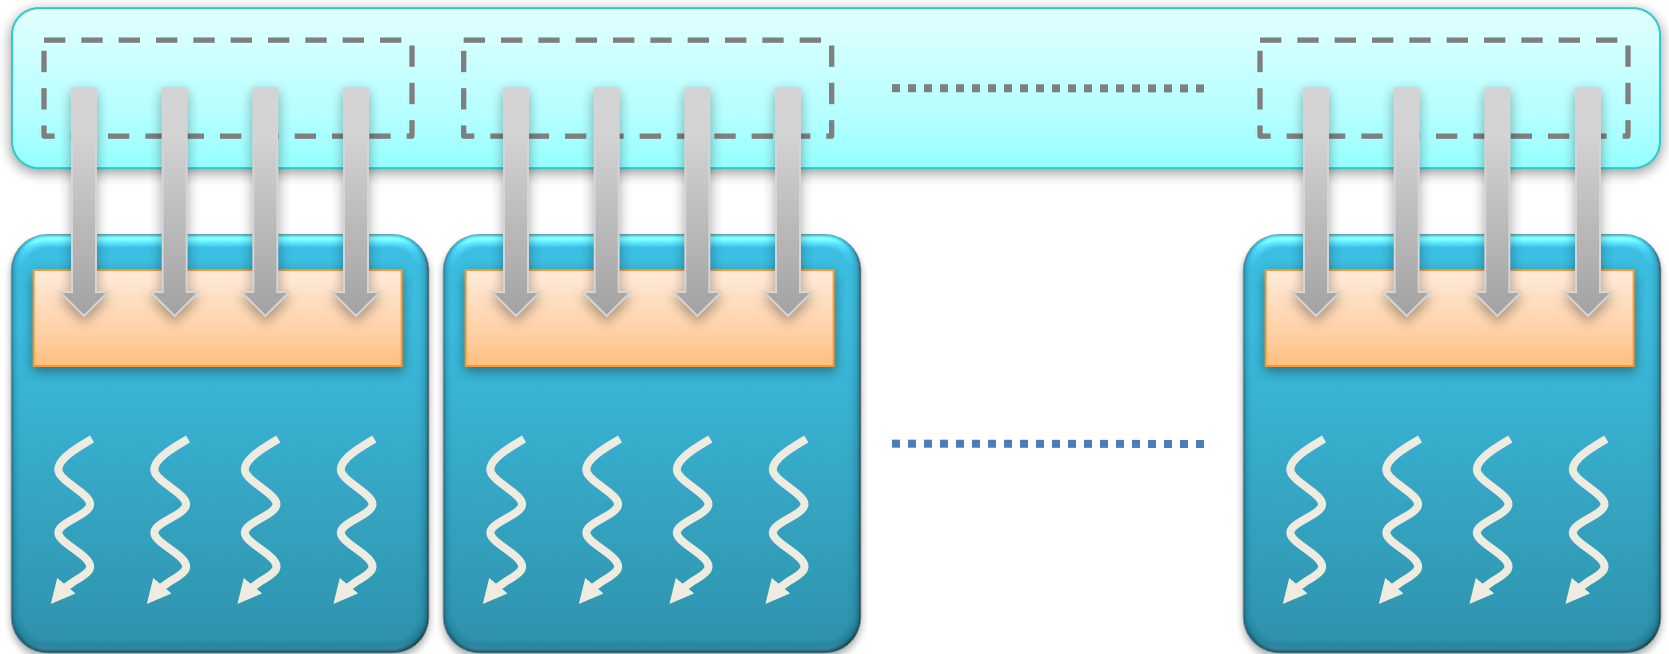
- Partition data into subsets that fit into shared memory

Primordial CUDA Pattern: Blocking



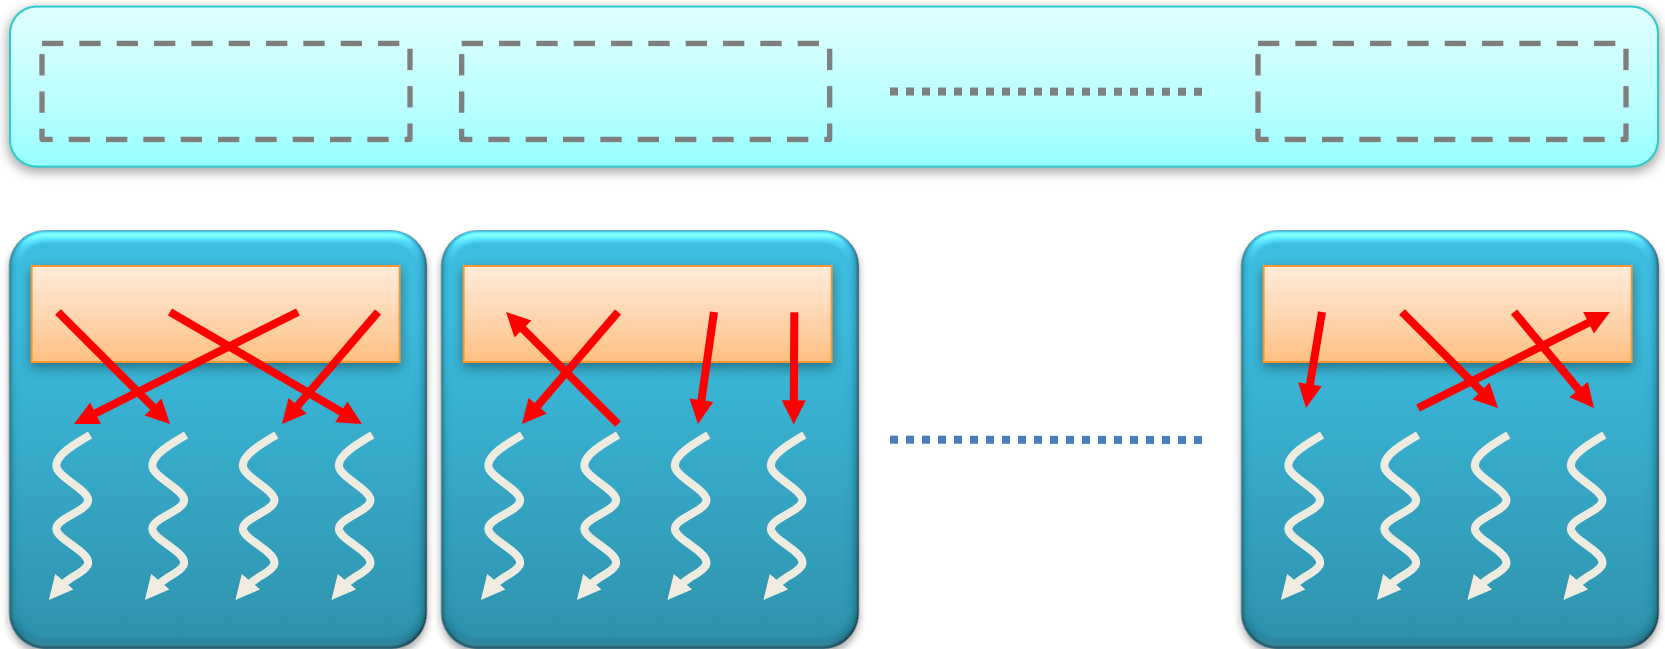
- Handle each data subset with one **thread block**

Primordial CUDA Pattern: Blocking



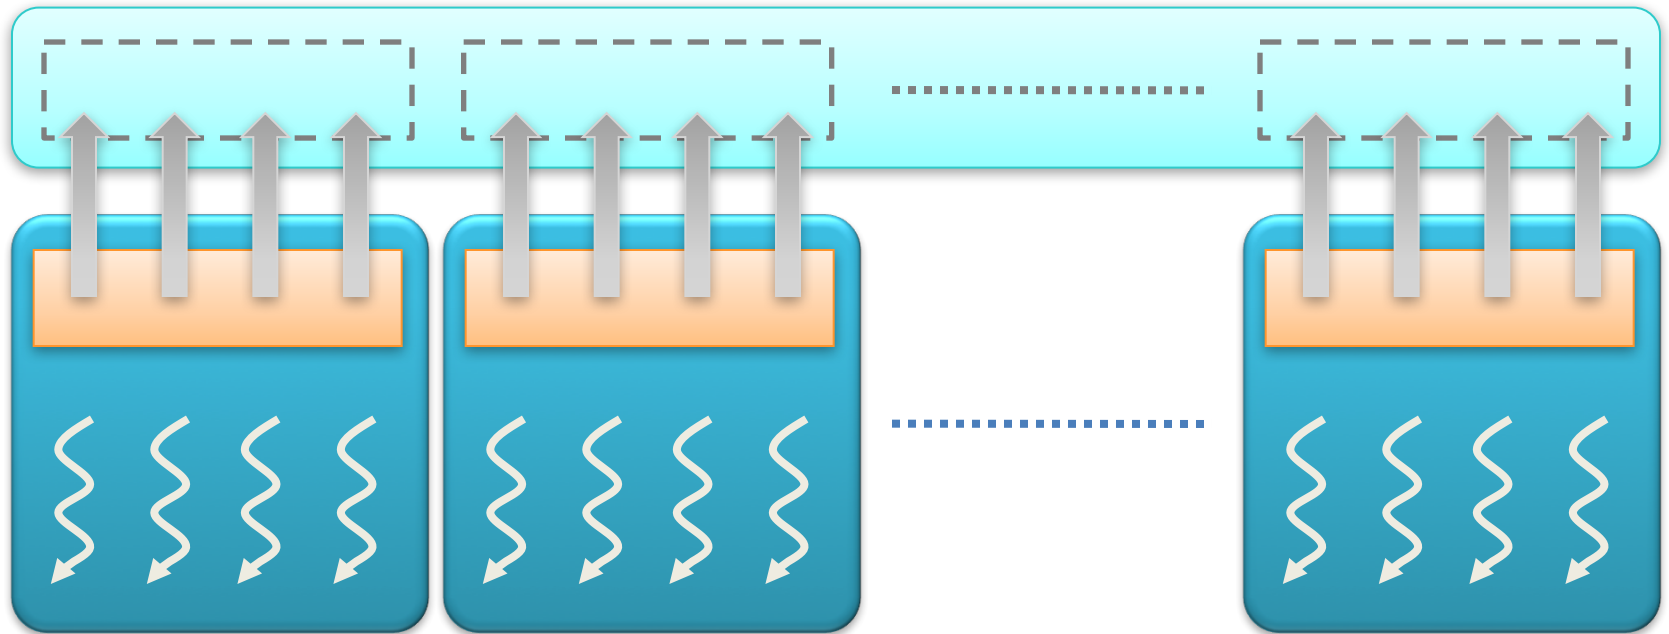
- Load the subset from global memory to shared memory, using multiple threads to exploit memory-level parallelism

Primordial CUDA Pattern: Blocking



- Perform the computation on the subset from **shared memory**

Primordial CUDA Pattern: Blocking



- Copy the result from **shared memory** back to global memory

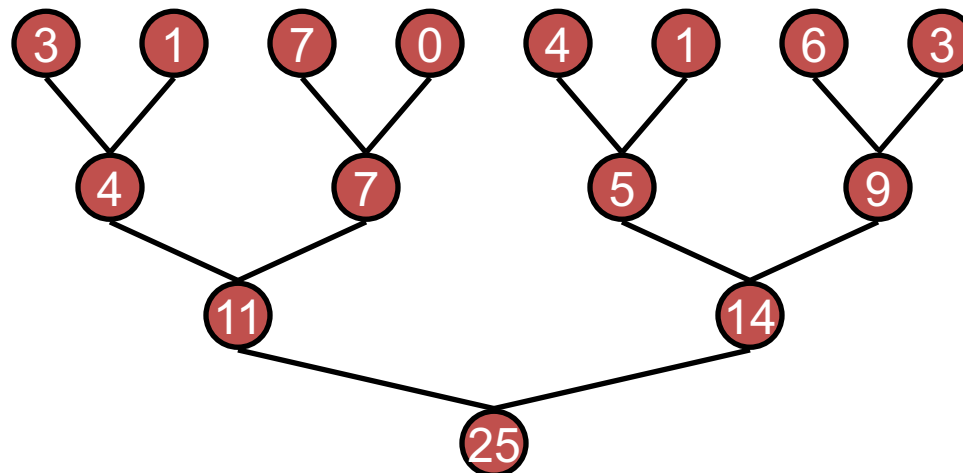
Primordial CUDA Pattern: Blocking

- All CUDA kernels are built this way
 - Blocking may not matter for a particular problem, but you're still forced to think about it
 - Not all kernels require shared memory
 - All kernels do require registers
- All of the parallel patterns we'll discuss have CUDA implementations that exploit blocking in some fashion

Reduction

- **Reduce** vector to a single value
 - Via an associative operator (+, *, min/max, AND/OR, ...)
 - CPU: sequential implementation 교환법칙이 성립하는 연산자

```
for(int i = 0, i < n, ++i) ...
```
 - GPU: “tree”-based implementation



Serial Reduction

```
// reduction via serial iteration
float sum(float *data, int n) {
    float result = 0;
    for(int i = 0; i < n; ++i)
    {
        result += data[i];
    }

    return result;
}
```

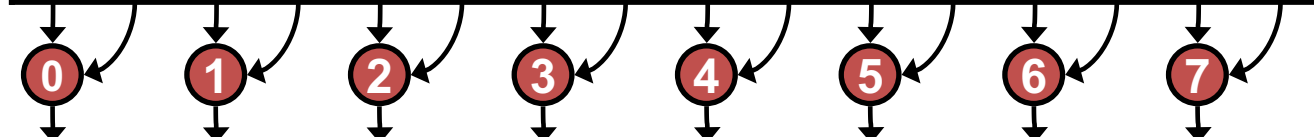
Parallel Reduction – Interleaved

Values (in shared memory)

Step 1
Stride 1

Thread
IDs

10	1	8	-1	0	-2	3	5	-2	-3	2	7	0	11	0	2
----	---	---	----	---	----	---	---	----	----	---	---	---	----	---	---

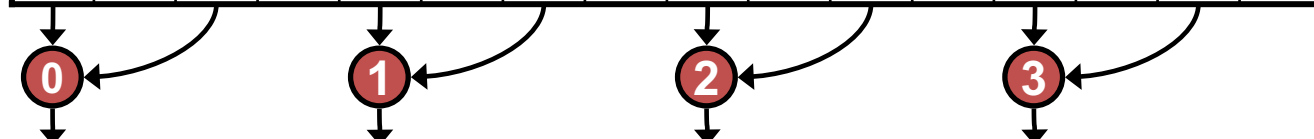


Values

11	1	7	-1	-2	-2	8	5	-5	-3	9	7	11	11	2	2
----	---	---	----	----	----	---	---	----	----	---	---	----	----	---	---

Step 2
Stride 2

Thread
IDs

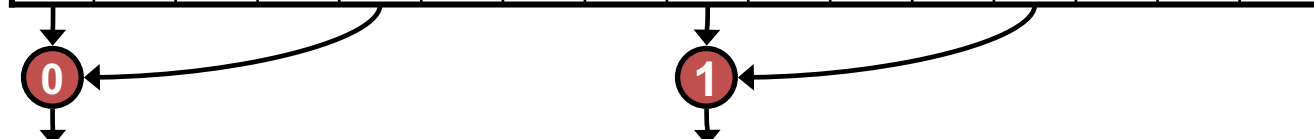


Values

18	1	7	-1	6	-2	8	5	4	-3	9	7	13	11	2	2
----	---	---	----	---	----	---	---	---	----	---	---	----	----	---	---

Step 3
Stride 4

Thread
IDs



Values

24	1	7	-1	6	-2	8	5	17	-3	9	7	13	11	2	2
----	---	---	----	---	----	---	---	----	----	---	---	----	----	---	---

Step 4
Stride 8

Thread
IDs



Values

41	1	7	-1	6	-2	8	5	17	-3	9	7	13	11	2	2
----	---	---	----	---	----	---	---	----	----	---	---	----	----	---	---

Parallel Reduction – Contiguous

Values (in shared memory)

Step 1
Stride 8

Thread
IDs

10	1	8	-1	0	-2	3	5	-2	-3	2	7	0	11	0	2
----	---	---	----	---	----	---	---	----	----	---	---	---	----	---	---



Values

8	-2	10	6	0	9	3	7	-2	-3	2	7	0	11	0	2
---	----	----	---	---	---	---	---	----	----	---	---	---	----	---	---

Step 2
Stride 4

Thread
IDs



Values

8	7	13	13	0	9	3	7	-2	-3	2	7	0	11	0	2
---	---	----	----	---	---	---	---	----	----	---	---	---	----	---	---

Step 3
Stride 2

Thread
IDs



Values

21	20	13	13	0	9	3	7	-2	-3	2	7	0	11	0	2
----	----	----	----	---	---	---	---	----	----	---	---	---	----	---	---

Step 4
Stride 1

Thread
IDs



Values

41	20	13	13	0	9	3	7	-2	-3	2	7	0	11	0	2
----	----	----	----	---	---	---	---	----	----	---	---	---	----	---	---

CUDA Reduction

```
__global__ void block_sum(float *input,
                          float *results,
                          size_t n) {
    extern __shared__ float sdata[];
    int i = ..., int tx = threadIdx.x;

    // load input into __shared__ memory
    float x = 0;
    if(i < n)
        x = input[i];
    sdata[tx] = x;
    __syncthreads();
}
```


CUDA Reduction

```
// block-wide reduction in __shared__ mem
for(int offset = blockDim.x / 2;
    offset > 0;
    offset >>= 1)
{
    if(tx < offset)
    {
        // add a partial sum upstream to our own
        sdata[tx] += sdata[tx + offset];
    }
    __syncthreads();
}
```

CUDA Reduction

```
// finally, thread 0 writes the result
if (threadIdx.x == 0)
{
    // note that the result is per-block
    // not per-thread
    results[blockIdx.x] = sdata[0];
}
}
```

An Aside

```
// is this barrier divergent?  
for(int offset = blockDim.x / 2;  
    offset > 0;  
    offset >>= 1)  
{  
    ...  
    __syncthreads();  
}
```

An Aside

```
// what about this one?
```

```
__global__ void do_i_halt(int *input)
{
    int i = ...
    if(input[i])
    {
        ...
        __syncthreads(); // a divergent barrier
                          // hangs the machine
    }
}
```

CUDA Reduction

```
// global sum via per-block reductions
float sum(float *d_input, size_t n)
{
    size_t block_size = ..., num_blocks = ...;

    // allocate per-block partial sums
    // plus a final total sum
    float *d_sums = 0;
    cudaMalloc((void**) &d_sums,
        sizeof(float) * (num_blocks + 1));
    ...
}
```

CUDA Reduction

```
// reduce per-block partial sums
int smem_sz = block_size*sizeof(float);
block_sum<<<num_blocks,block_size,smem_sz>>>(d_input, d_sums, n);

// reduce partial sums to a total sum
block_sum<<<1,block_size,smem_sz>>> (d_sums, d_sums + num_blocks, num_blocks);

// copy result to host
float result = 0;
cudaMemcpy(&result, d_sums+num_blocks, ...);
return result;
}
```

Caveat Reductor

- What happens if there are too many partial sums to fit into shared memory in the second stage?
- What happens if the temporary storage is too big?
- Give each thread more work in the first stage
 - Sum is **associative** & **commutative**
 - Order doesn't matter to the result
 - We can schedule the sum any way we want
 - serial accumulation before block-wide reduction
- Exercise left to the students

Parallel Reduction Complexity

- $\text{Log}(N)$ parallel steps, each step S does $N/2^S$ independent ops
 - **Step Complexity** is $O(\log N)$
- For $N=2^D$, performs $\sum_{S \in [1..D]} 2^{D-S} = N-1$ operations
 - **Work Complexity** is $O(N)$ – It is **work-efficient**
 - i.e. does not perform more operations than a sequential algorithm
- With P threads physically in parallel (P processors), **time complexity** is $O(N/P + \log N)$
 - Compare to $O(N)$ for sequential reduction

Common Parallel Computing Scenarios

- Many parallel threads need to generate a single result
→ Reduce
- Many parallel threads need to partition data
→ Split
- Many parallel threads produce variable output / thread
→ Compact / Expand

Split Operation

- Given: array of true and false elements (and payloads)

Flag

T	F	F	T	F	F	T	F
---	---	---	---	---	---	---	---

Payload

3	1	7	0	4	1	6	3
---	---	---	---	---	---	---	---

- Return an array with all true elements at the beginning

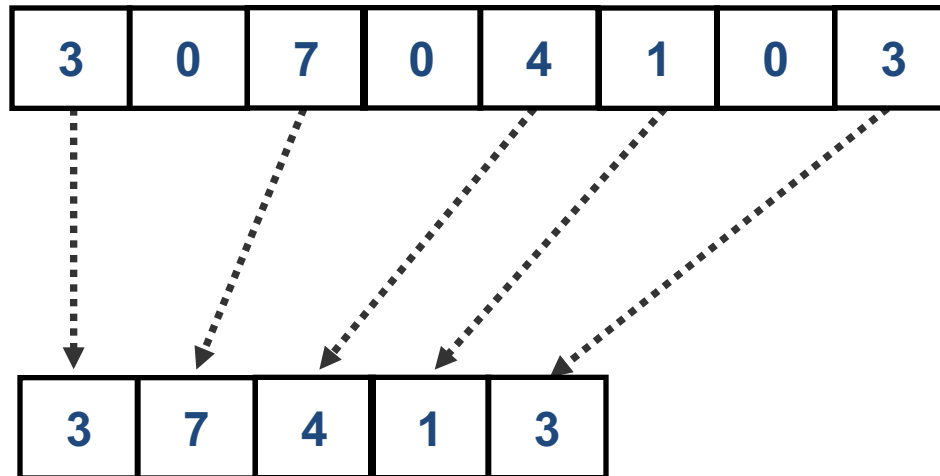
T	T	T	F	F	F	F	F
---	---	---	---	---	---	---	---

3	0	6	1	7	4	1	3
---	---	---	---	---	---	---	---

- Examples: sorting, building trees

Variable Output Per Thread: Compact

- Remove null elements



Variable Output Per Thread: General Case

- Reserve Variable Storage Per Thread

2	1	0	3	2
---	---	---	---	---

A	C
B	

D	G
E	H
F	

- Example: binning

Split, Compact, Expand

- Each thread must answer a simple question:

“Where do I write my output?”

- The answer depends on what other threads write!
- *Scan* provides an efficient parallel answer

Scan (a.k.a. Parallel Prefix Sum)

- Given an array $A = [a_0, a_1, \dots, a_{n-1}]$ and a binary associative operator \oplus with identity I ,

$$\text{scan}(A) = [I, a_0, (a_0 \oplus a_1), \dots, (a_0 \oplus a_1 \oplus \dots \oplus a_{n-2})]$$

- Prefix sum: if \oplus is addition, then scan on the series

3	1	7	0	4	1	6	3
---	---	---	---	---	---	---	---

returns the series

0	3	4	11	11	15	16	22
---	---	---	----	----	----	----	----

Applications of Scan

- Scan is a simple and useful parallel building block for many parallel algorithms:
 - Radix sort
 - Quicksort (seg. scan)
 - String comparison
 - Lexical analysis
 - Stream compaction
 - Run-length encoding
 - Polynomial evaluation
 - Solving recurrences
 - Tree operations
 - Histograms
 - Allocation
 - Etc.
- Fascinating, since scan is unnecessary in sequential computing!

Serial Scan

```
int input[8] = {3, 1, 7, 0, 4, 1, 6, 3};
int result[8];
int running_sum = 0;
for(int i = 0; i < 8; ++i)
{
    result[i] = running_sum;
    running_sum += input[i];
}

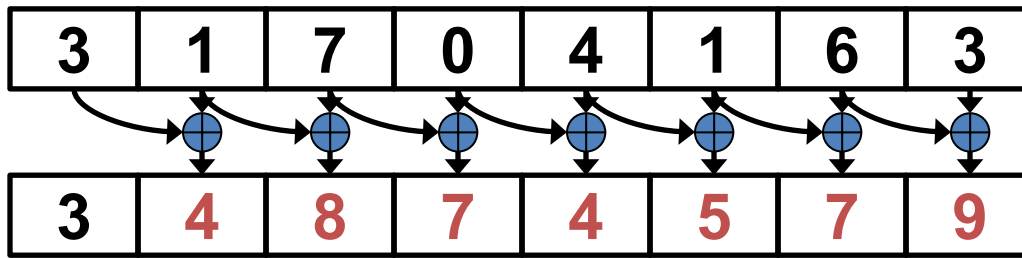
// result = {0, 3, 4, 11, 11, 15, 16, 22}
```


A Scan Algorithm – Preview


3	1	7	0	4	1	6	3
---	---	---	---	---	---	---	---

Assume array is already in shared memory

A Scan Algorithm – Preview

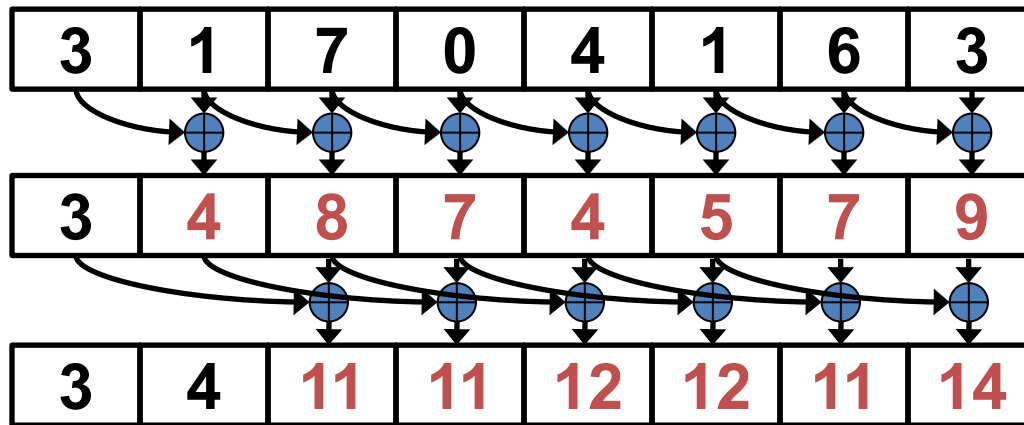


Iteration 0, $n-1$ threads


Each  corresponds to a single thread.

Iterate $\log(n)$ times. Each thread adds value *stride* elements away to its own value

A Scan Algorithm – Preview

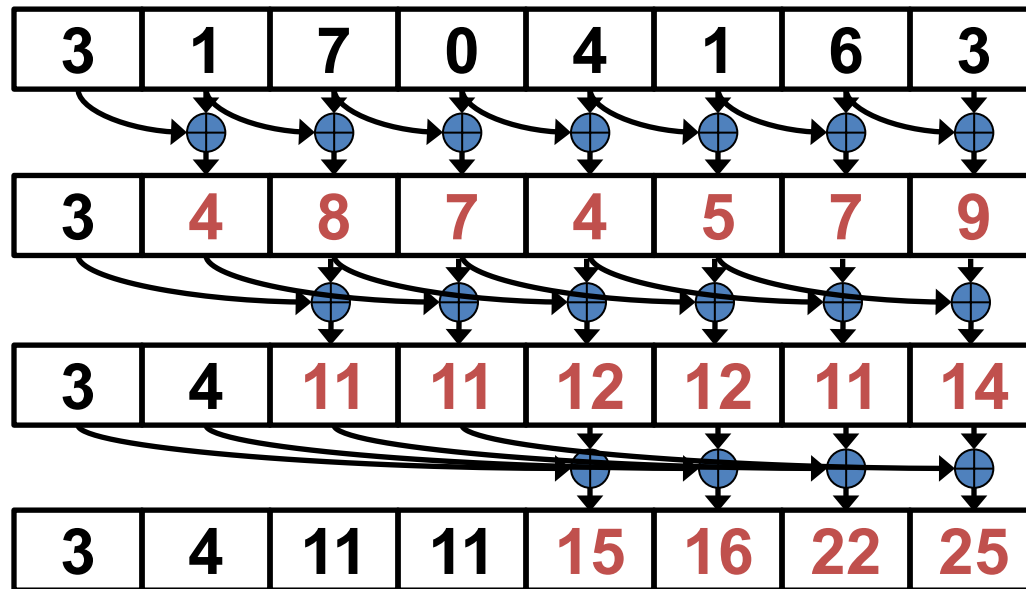


Iteration 1, $n-2$ threads


Each  corresponds to a single thread.

Iterate $\log(n)$ times. Each thread adds value *offset* elements away to its own value

A Scan Algorithm – Preview



Iteration i , $n-2^i$ threads

Each  corresponds to a single thread.

Iterate $\log(n)$ times. Each thread adds value *offset* elements away to its own value.

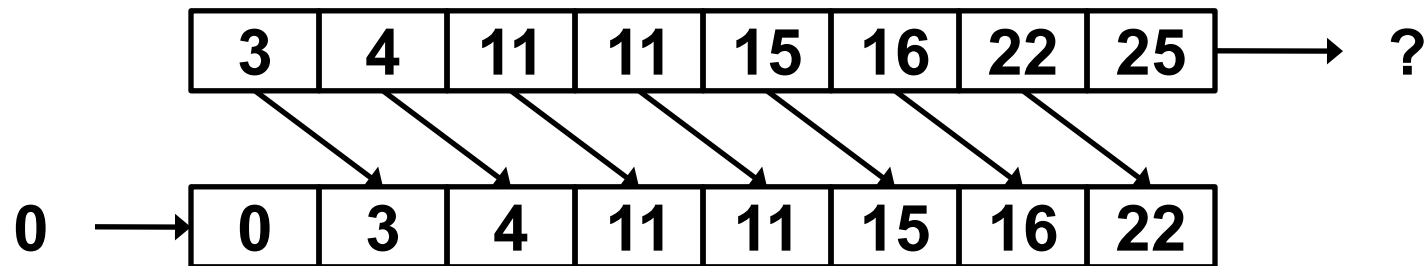
Note that this algorithm operates in-place: no need for double buffering

A Scan Algorithm – Preview

3	4	11	11	15	16	22	25
---	---	----	----	----	----	----	----

- We have an **inclusive** scan result

A Scan Algorithm – Preview



- For an **exclusive** scan, right-shift through **__shared__** memory
- Note that the unused final element is also the sum of the entire array
 - Often called the “carry”
 - Scan & reduce in one pass

CUDA Block-wise Inclusive Scan

```
__global__ void inclusive_scan(int *data)
{
    extern __shared__ int sdata[];

    unsigned int i = ...

    // load input into __shared__ memory
    int sum = input[i];
    sdata[threadIdx.x] = sum;
    __syncthreads();
    ...
}
```

CUDA Block-wise Inclusive Scan

```
for(int o = 1; o < blockDim.x; o <<= 1) {  
    if(threadIdx.x >= o)  
        sum += sdata[threadIdx.x - o];  
    // wait on reads  
    __syncthreads();  
  
    // write my partial sum  
    sdata[threadIdx.x] = sum;  
  
    // wait on writes  
    __syncthreads();  
}
```


CUDA Block-wise Inclusive Scan

```
// we're done!  
// each thread writes out its result  
result[i] = sdata[threadIdx.x];  
}
```

Results are Local to Each Block

Block 0

Input:

5 5 4 4 5 4 0 0 4 2 5 5 1 3 1 5

Result:

5 10 14 18 23 27 27 27 31 33 38 43 44 47 48 53

Block 1

Input:

1 2 3 0 3 0 2 3 4 4 3 2 2 5 5 0

Result:

1 3 6 6 9 9 11 14 18 22 25 27 29 34 39 39

Results are Local to Each Block

- Need to propagate results from each block to all subsequent blocks
- 2-phase scan
 1. Per-block scan & reduce
 2. Scan per-block sums
- Final update propagates phase 2 data and transforms to exclusive scan result

Summing Up

- Patterns like **reduce**, **split**, **compact**, **scan**, and others let us reason about data parallel problems abstractly
- Higher level patterns are built from more fundamental patterns
- Scan in particular is **fundamental** to parallel processing, but unnecessary in a serial world